

Context-Augmented, Session-Calibrated Blood Pressure Tracking from Photoplethysmography

Farouk Ganiyu Adewumi and Timothy Oladunni

Department of Computer Science, Morgan State University, Baltimore, MD 21251, USA

fagan1@morgan.edu; timothy.oladunni@morgan.edu

Abstract

Objective. To determine whether acquisition context and photoplethysmography (PPG) features improve session-calibrated blood pressure tracking while controlling participant leakage and calibration burden.

Approach. We performed repeated five-fold participant-level cross-validation on 653 AuroraBP participants, excluding calibration-labelled rows from fitting and scoring. We also conducted a feature-family ablation, an independent 108-participant pulse-rate fusion analysis, and a separately developed VitalDB evaluation on 227 held-out subjects.

Main results. With a three-reading cuff anchor, the wearable-state model achieved participant-level mean absolute errors of 7.697 mmHg for systolic blood pressure (SBP) and 6.139 mmHg for diastolic blood pressure (DBP). Wearable context reduced error by 0.436 mmHg for SBP (97.5% confidence interval 0.305–0.573) and 0.126 mmHg for DBP (0.053–0.200). Standard PPG descriptors provided further reductions of 0.187 mmHg for SBP and 0.180 mmHg for DBP, whereas nonlinear/attractor features did not show a reliable benefit. Quality-weighted visible/infrared fusion reduced pulse-rate error from 3.985 to 3.690 beats per minute. In VitalDB, combining context and PPG predictions reduced error by 0.451 mmHg for SBP and 0.357 mmHg for DBP, although absolute error remained 15.375 and 8.833 mmHg, respectively.

Significance. Context and standard PPG descriptors produced reproducible but clinically small improvements. The results support calibrated tracking research but do not establish model transfer, equitable performance, standards compliance, or cuff replacement.

Keywords: photoplethysmography; cuffless blood pressure; session calibration; participant-level cross-validation; wearable sensing; data leakage; signal quality.

1 Introduction

1.1 Motivation

PPG-based cuffless BP estimation is attractive because optical sensors are already present in consumer wearables. Its evaluation must distinguish participant overlap, calibration burden, and acquisition-context shortcuts. Record-level splits can inflate apparent accuracy when correlated measurements from one participant enter both training and testing (González et al., 2023). A model anchored to three initial cuff readings also answers a different question from one calibrated once or repeatedly supplied with recent cuff values. Finally, pigmentation-associated optical differences motivate intervention testing, but a small subgroup or a nonsignificant contrast cannot establish equitable performance (Kostrna et al., 2026; Singh et al., 2024; Osorio-Sanchez et al., 2025). We audit anchoring, context, and pigmentation-related interventions using recorded Fitzpatrick labels. Candidate interventions and context arms were frozen before repeated participant-level validation, after earlier development on AuroraBP. Because a fixed holdout was inspected across successive ablations, it is retained as historical development evidence. The primary repeated analysis confines feature ranking and fitting to outer training participants but remains internal validation. Secondary negative results are reported without interpreting them as proof of equity, equivalence, or general failure of pigmentation-aware modeling.

1.2 Study Questions and Contributions

The primary question is whether wearable-state context improves session-anchored SBP and DBP tracking beyond the same calibrated model core. Secondary analyses test calibration dose, standard versus nonlinear/attractor PPG feature families, and pigmentation-related interventions. Historical development analyses examine how recent-cuff anchoring can inflate apparent gains. These questions were specified retrospectively; repeated cross-validation estimates internal stability rather than prospective confirmation.

The study contributes participant-separated evaluation of wearable context, a calibration-row leakage audit, and matched calibration-dose and feature-family comparisons. Independent analyses evaluate visible/infrared pulse-rate fusion and a separately developed VitalDB context-plus-PPG model. Negative pigmentation-intervention and attractor-ablation findings are retained without claims of equity or physiological validation. Historical development results and auxiliary diagnostics appear in the supplementary material.

2 Related Work

2.1 PPG-Based Cuffless Blood Pressure Estimation

Cuffless BP estimation from PPG includes pulse-transit-time methods, waveform features, and end-to-end networks (Elgendi et al., 2019; Mukkamala et al., 2015). Calibration-free methods avoid subject-specific references, whereas calibrated approaches use one or more cuff readings to model subsequent change (Samimi and Dajani, 2023; Wang et al., 2023; Oladunni and Adewumi, 2026b). Evaluation guidance emphasizes participant separation, calibration burden, and careful use of AAMI/ISO 81060-2 and BHS criteria (Elgendi et al., 2024).

2.2 Nonlinear and Information-Theoretic Feature Representations

The feature bank combines waveform moments with delay-coordinate reconstruction, recurrence analysis, sample entropy, permutation entropy, and spectral measures (Takens, 1981; Eckmann et al., 1987; Marwan et al., 2007; Richman and Moorman, 2000; Bandt and Pompe, 2002). These established descriptors have been applied to physiological and arterial pulse signals (Allen, 2007; Hörandtner et al., 2022). Recent author-developed formulations motivated the feature grouping and terminology (Oladunni and Adewumi, 2026c,b,a), but are not treated as validated physiological laws. The matched ablation therefore tests the empirical contribution of the combined nonlinear/attractor family rather than using prediction as evidence for the underlying theories.

2.3 Skin Pigmentation and PPG-Derived Vitals

Pigmentation can affect optical measurements, motivating separate evaluation of signal quality and downstream error (Singh et al., 2024; Kostrna et al., 2026; Osorio-Sanchez et al., 2025; Al-Halawani et al., 2025). Existing field, polarization, and gain-correction studies support testing optical interventions but do not establish improved BP tracking (Dasa and Davies, 2026; Jakachira et al., 2026; MacIsaac et al., 2025). We therefore distinguish pulse extraction from BP prediction and do not interpret nonsignificant subgroup differences as parity.

2.4 Evaluation Rigor and Data Leakage

AuroraBP established calibration-based benchmarks and documented differences between laboratory and ambulatory performance (Mieloszyk et al., 2022). Record- or window-level

splits can inflate reported accuracy when one participant contributes to training and testing (González et al., 2023; Sanches et al., 2024), while dataset shortcuts may further obscure the information contributed by PPG (Dias et al., 2025; Mehta et al., 2024). Our audit addresses a related change-prediction artifact: a calibration-row indicator may reveal observations whose target is near zero even when participant folds are disjoint.

3 Framework and Audit Controls

3.1 Attractor Feature Bank

Let x_n , $n = 1, \dots, T$, denote a discretely sampled PPG window at rate f_s . Following delay-coordinate embedding (Takens, 1981), we reconstruct the trajectory

$$\mathbf{X} = \tau_{\mathcal{G}}[x], \quad \mathbf{X}_{n,\cdot} = [x_n, x_{n+\tau}, \dots, x_{n+(d-1)\tau}], \quad (1)$$

with embedding dimension $d = 4$ and delay $\tau = 5$ samples, matching the AVCT/CST instantiation (Oladunni and Adewumi, 2026b,c). From $\mathbf{X}$ we compute the shared PPG attractor feature bank

$$\Phi(x) = [\sigma_m, \text{skew}_m, \text{kurt}_m, H_{\text{SE}}, H_{\text{PE}}, R_{\text{rec}}, R_{\text{det}}, H_{\text{RQA}}, \text{CSI}, H_{\text{spec}}, Q] \quad (2)$$

comprising waveform moments (σ_m , skewness, kurtosis), sample and permutation entropy (H_{SE} , H_{PE}), recurrence rate and determinism (R_{rec} , R_{det}), recurrence entropy (H_{RQA}), CSI, spectral entropy, and optical quality. The auxiliary stress-test bank labels the unmodified measured waveform **ref_***, the simulated stressed waveform **raw_***, and its corrected version **pc_***; “reference” does not mean latent physiological ground truth. AuroraBP instead extracts **raw_*** from observed optical waveforms and **pc_*** after the correction below. It does not construct an AuroraBP **ref_*** input. The primary candidate pool contains raw-feature changes, time variables, and a UCI-derived prior; its exact provenance is described in Section 4.2.

3.2 Pigmentation Correction Operator

The implemented AuroraBP operator is a heuristic motivated by attenuation, rather than an estimate of melanin or a fitted Beer–Lambert law. Let $z = (\text{FST} - 1)/5$, $b = \text{median}(x)$, $a = \max(0.45, 1 - 0.55z)$, and $u = (x - b)/a$. With $W(u)$ a Wiener-filtered version of u and $\lambda = \text{clip}(0.75z, 0, 0.75)$, the corrected channel is $b + (1 - \lambda)u + \lambda W(u)$. The filter kernel

is an odd number of samples derived from $0.05f_s$, bounded between 5 and 21. A residual-to-bandpassed-signal SD ratio, clipped to $[0, 1.5]$, modifies entropy/recurrence tolerances; the corrected channel uses 0.35 times that ratio. The waveform median is a baseline proxy, not a separate measured DC channel. In the repeated participant-CV ablation, corrected-minus-raw MAE was -0.007 mmHg SBP (95% CI -0.042 to 0.029) and -0.030 mmHg DBP (-0.062 to 0.003). These intervals do not establish overall benefit or equivalence. The ablation tests the complete correction-and-feature-extraction configuration, not attenuation rescaling alone.

3.3 Calibration Dose and Anchoring

For participant p , let $b_p^{(k)}$ denote an anchor formed from exactly k initial cuff readings ($k \in \{1, 2, 3\}$), with $k > 1$ using their arithmetic mean. We define the *session-anchored* target

$$\Delta_{\text{session}}^{(k)}(p, i) = y_{p,i} - b_p^{(k)}, \quad (3)$$

where $y_{p,i}$ is the i -th subsequent cuff reading for participant p , and the *prior-anchored* target

$$\Delta_{\text{prior}}(p, i) = y_{p,i} - y_{p,i-1}^*, \quad (4)$$

where $y_{p,i-1}^*$ is the most recent valid prior measurement (which may itself be a tracked, non-calibration reading). **These are not the same question.** Eqn. 3 evaluates BP change after an initial calibration session without supplying later cuff values to the predictor. Eqn. 4 asks whether PPG can track change from the most recent cuff reading, which in this in-lab protocol can be only minutes old. A model evaluated under Eqn. 4 solves an easier problem; 83% of the historical fixed-split gain under prior anchoring was attributable to anchor replacement alone. We therefore lead with session-anchored tracking results and explicitly report calibration dose; neither anchoring scheme establishes cuff replacement. “One reading” refers only to $k = 1$; the primary $k = 3$ analysis is described as *one-session calibration using three readings*. The *calibration-retention baseline* predicts $\hat{\Delta} = 0$ under either anchor, i.e., “the BP has not changed since calibration (or since the last reading).” This is one baseline; the main context contrasts compare fitted models, and VitalDB additionally uses a calibration-plus-context control.

3.4 Context Operator

Let $c(p, i)$ collect acquisition context. We distinguish a *wearable-state* set (posture, activity, and optical signal quality) from a *protocol-enriched* set that additionally includes protocol

phase, recording duration, and pressure-quality metadata. The wearable-state variables require corresponding sensing or classification at deployment; the protocol-enriched variables may encode AuroraBP acquisition structure and are therefore interpreted as an internal upper bound on context utility. The context-augmented model is

$$\hat{\Delta}(p, i) = f_{\theta}(\Phi(x_{p,i}), \Phi(x_{p,i}) - \Phi(x_{p,\text{anchor}}), c(p, i)), \quad (5)$$

where f_{θ} is a gradient-boosted tree ensemble (HistGradientBoosting; hyperparameters grid-searched over learning rate, maximum leaf nodes, and ℓ_2 regularization) trained with either squared or absolute loss (Section 4), and the first-difference term $\Phi(x_{p,i}) - \Phi(x_{p,\text{anchor}})$ is the attractor-feature analogue of Eqns. 3–4 in feature space.

3.5 Calibration-Row Leakage Audit

A calibration indicator could identify trivial rows if a target reading were compared with itself, allowing a model to learn “calibration row $\Rightarrow$ predict no change.” We used four checks:

- A** Compare $|\Delta|$ on flagged and unflagged rows.
- B** Re-score the same fitted models with flagged rows removed from the holdout only.
- C** Refit with the calibration-flag feature removed, flagged rows still scored.
- D (Decisive)** Remove the flag and exclude flagged rows from fitting and scoring.

Checks A–C each leave one channel open; Check D restricts both fitting and evaluation. Supplementary Section B-F reports all four.

Finding 1. *On this dataset, calibration-flagged rows are not near-zero-delta (mean $|\Delta| = 6.774$ vs. 9.902 mmHg SBP on development rows – smaller, but not trivially so), and Check D confirms the context gain survives essentially intact ($+0.614$ vs. the unaudited $+0.809$ mmHg SBP). Leave-one-out results implicate recording duration and post-exercise activity rather than the calibration flag (Supplementary Section B-F).*

4 Methodology

4.1 Datasets

AuroraBP is the primary dataset (Mieloszyk et al., 2022; Microsoft Research, 2022). It contains 1,221 participant records and 11,597 auscultatory measurements across calibration,

static, exercise, and temporal-challenge phases; waveforms were resolvable for 672 participants. The primary common cohort includes 653 participants with three initial calibration readings. Excluding 1,337 later calibration-labelled rows leaves 7,879 scored measurements. Historical ablations used a fixed 501/168 split; because its holdout was inspected during development, it is treated as historical evidence rather than untouched validation. Four secondary sources appear in the development notebooks: BIDMC, MIMIC-II, Welltory-PPG, and UCI (Pimentel et al., 2018, 2017; Goldberger et al., 2000; Saeed et al., 2011b,a; Neshitov et al., 2021; Kachuee et al., 2015b,a). The first three form a historical corrected-feature PCA bank whose coordinates are excluded from the primary candidate list. UCI trains the BP-change prior that may enter the primary model. Pigmentation ranks on these sources are synthetic simulation inputs, not measured traits, and support no equity claim.

Two datasets support independent follow-up analyses. The **Appel–Theart Participant Study Data** provide paired visible and infrared PPG, Monk Skin Tone (MST), and pulse-rate references (Appel and Theart, 2025). After auditing, 108 participants contributed 630 pulse-rate windows, including 24 participants with MST 7–10. Its single cuff BP value per participant permits only cross-sectional BP analysis. **VitalDB** provides synchronized intraoperative PPG and arterial BP (Lee et al., 2022; Lee and Jung, 2022). A locked evaluation used 5,472 future windows from 228 cases belonging to 227 subjects disjoint from the 55 development subjects. Cases were grouped by subject for inference. VitalDB has no skin-tone field and tests a separately developed intraoperative implementation, not transfer of AuroraBP-fitted parameters. Table 1 consolidates the analysis sample, role, and source references for every dataset used in the study.

4.2 Feature Extraction and Selection

AuroraBP features are extracted from the central 10-s optical window at $f_s = 500$ Hz (5,000 samples), using a third-order 0.5–8 Hz bandpass, $d = 4$, and $\tau = 5$ samples (10 ms). The extractor calculates $f_s = (N - 1)/(t_{N-1} - t_0)$ from the recording timestamps; all 11,427 cached feature rows match the timestamp audit at 500 Hz within 2×10^{-11} Hz. The central waveform is not resampled before filtering. Sample entropy uses order 2, a base standardized tolerance of 0.2, and at most 600 samples; permutation entropy uses order 3. Recurrence analysis uses at most 300 standardized embedded points, Chebyshev distance, a base radius of 0.2 modified by the noise ratio, and diagonal lines of length at least 2. For the filtered AuroraBP signal, let $S(f)$ be its Welch power spectral density, with 4-s Hann segments and

Table 1: Datasets Used

Dataset	Material analyzed	Role in this study	Source and access
AuroraBP	1,221 released participant records; 653 participants and 7,879 measurements in the strict primary endpoint	Primary session-calibrated BP evaluation, context arms, calibration dose, subgroup diagnostics, and feature-family ablation	Microsoft Research study and Zenodo dataset; available upon application under the provider’s Data Use Agreement (Mieloszyk et al., 2022; Microsoft Research, 2022)
BIDMC	53 subjects; 1,272 windows	Historical corrected-feature PCA representation; not a primary evaluation cohort	BIDMC PPG and Respiration Dataset, PhysioNet (Pimentel et al., 2018, 2017; Goldberger et al., 2000)
MIMIC-II	1,500 sampled windows	Historical corrected-feature PCA representation; not a primary evaluation cohort	MIMIC-II Clinical Database, PhysioNet (Saeed et al., 2011b,a; Goldberger et al., 2000)
Welltory-PPG	106 windows	Historical corrected-feature PCA representation; not a primary evaluation cohort	Published PPG signal collection (Neshitov et al., 2021)
UCI Cuff-Less BP	425 records; 3,268 windows	Training source for the BP-change prior that may enter AuroraBP models	UCI Machine Learning Repository (Kachuee et al., 2015b,a)
Appel–Theart	108 analyzed participants; 630 pulse-rate windows	Independent visible/infrared quality-fusion and guarded MST-correction analysis	Mendeley Data (Appel and Theart, 2025)
VitalDB	55 development subjects; 228 held-out cases from 227 subjects; 5,472 scored future windows	Independent intraoperative calibration-plus-context and context-plus-PPG BP tracking	VitalDB and PhysioNet (Lee et al., 2022; Lee and Jung, 2022)

Counts describe the material used by this study and need not equal the full source repository. Auxiliary-source pigmentation ranks were synthetic simulation inputs; only AuroraBP supplied Fitzpatrick labels and Appel–Theart supplied Monk Skin Tone labels.

50% overlap. The quality feature is

$$Q = \frac{\sum_{0.5 \leq f \leq 5} S(f)}{\sum_{0.1 \leq f \leq 15} S(f)}. \quad (6)$$

CSI is $0.4R_{\text{det}} + 0.3[1 - \text{clip}(H_{\text{RQA}}/5, 0, 1)] + 0.3Q$. Recalculation from the cached components agrees with stored CSI to 2.3×10^{-16} . This derived spectral ratio, stored as `raw_ppg_quality`, is distinct from AuroraBP’s supplied `optical_quality` context field and the independent pulse-fusion weights q_V, q_I . The historical hybrid calibration-anchored feature augmentation is denoted HCA below; its source-code identifiers retain the older name H1, distinct from hypothesis H1.

Within each outer training fold, the primary code ranks raw-feature differences and relative differences, elapsed-time variables, and the UCI prior after ridge residualization against baseline BP and log elapsed time. It applies `mutual_info_regression` to the residuals and retains up to 24 candidates plus baseline BP. SMD and AUC were historical diagnostics, not the primary selection algorithm. The name “no context” means no *added* context variables: baseline BP, candidate elapsed-time terms, and the UCI prior can still enter its shared core.

No `ref_*` feature enters that core. The UCI prior is generated from corrected features, so absence of explicit Fitzpatrick terms does not prove absence of indirect Fitzpatrick dependence. The regenerated selection manifests retain the UCI prior in 26 of 30 endpoint/fold core models; all context arms share the corresponding selected core. Historical SD ratios are reported in Supplementary Table S4; they are not variance fractions.

4.3 Model and Training

For the primary AuroraBP comparisons, the supplied implementation uses median imputation followed by HistGradientBoosting with squared-error loss, learning rate 0.03, 15 maximum leaf nodes, ℓ_2 regularization 1.0, 300 iterations, and random seed 42. These settings are inherited from development rather than reselected using outer test folds. Row weights are inverse participant measurement counts, clipped at the training-set 2nd/98th percentiles and normalized to mean one. Participant folds are stratified by Fitzpatrick band, with seeds $42+26200+r$ for zero-based repeat r . The change gate is selected using fitted training predictions, not an independent inner validation set; reported activation is 1.0. This training-stage gate selection is an additional development limitation. Exploratory endpoint specialization uses a separate nested procedure.

4.4 Evaluation Protocol

Primary results use three repeats of five-fold participant-stratified outer cross-validation over the common 653-participant cohort. Every participant appears in an outer test fold once per repeat. Feature ranking, imputation, model fitting, and gate selection are repeated using only the corresponding outer training participants; no record or window from an outer test participant enters training. The final feature specification and context arms were frozen before this repeated analysis. Calibration-labelled rows are excluded from both fitting and scoring. Per-participant MAE is averaged across repeats, and 95% CIs are obtained by participant bootstrap (10,000 resamples); context contrasts use paired per-participant differences. This is repeated *internal* validation because AuroraBP informed earlier model development, not an external or prospective test. Historical fixed-split analyses use participant bootstrap (4,000 resamples) and are labelled accordingly. BHS and AAMI/ISO results are retained as indicative historical context, not compliance claims.

4.4.1 Claim hierarchy and multiplicity

The two primary retrospective comparisons are SBP and DBP wearable-state context versus the same core without added context. The rerun saves each participant’s mean MAE across

the three repeats and forms paired no-context-minus-wearable-state contrasts. Two-sided paired sign-flip tests use the absolute mean contrast, 99,999 random sign assignments (seed 20260910), and $(b + 1)/(B + 1)$ Monte Carlo p values. Holm’s step-down procedure controls multiplicity across these two tests at $\alpha = 0.05$ (Holm, 1979). Paired t tests are reported as a sensitivity analysis. No p value is reconstructed from a CI width. Primary contrasts additionally report 97.5% participant-bootstrap intervals, giving approximate Bonferroni simultaneous 95% coverage for the two endpoints; other intervals are marginal 95%. The sign-flip test assumes sign exchangeability under the paired null. Inference conditions on the fitted out-of-fold predictions, treats participants as the sampling units, and does not capture all dependence from shared training sets or earlier model development. Multiplicity correction does not make this retrospectively designated family prospective confirmation. All other comparisons are secondary or exploratory. The matched feature-family ablation is a secondary analysis on the same strict cohort and outer folds. Every arm uses baseline BP, elapsed-time terms, wearable-state context, the same fixed estimator, participant weights, and training-stage gate; none uses the UCI prior. The standard-PPG arm adds changes and relative changes in amplitude dispersion, skewness, kurtosis, spectral entropy, the band-power quality ratio, and pulse-arrival time. The attractor-augmented arm additionally adds sample entropy, permutation entropy, recurrence rate, determinism, RQA entropy, and CSI. Fixed semantic feature sets replace foldwise ranking for this ablation so the contrast tests the declared family addition. Paired participant-bootstrap intervals are marginal; sign-flip tests for the two attractor-versus-standard endpoints receive a separate Holm correction. Because several nonlinear complexity measures are not uniquely attractor-derived, this is evidence about the combined nonlinear/attractor augmentation rather than any single descriptor. For the fixed-estimator pigmentation diagnostic, the prior-anchored reference model was refitted on the original 501 development participants and evaluated on the same 168 held-out participants. Fitzpatrick bands stratify the resulting predictions; they do not alter the fitted estimator or its gates. The upstream corrected-feature UCI prior can introduce indirect Fitzpatrick dependence. MAE is averaged within participant before band means, 10,000-resample bootstrap intervals, and exploratory between-band contrasts are computed. A separate band summary uses the regenerated primary wearable-state predictions from all 653 participants. These diagnostics cannot establish fairness, particularly with only 6 and 25 V–VI participants in the respective cohorts. The later endpoint-specialization analyses are explicitly exploratory. They use the same 653-participant strict cohort and three repeats of five-fold participant outer validation, with three-fold participant inner validation for candidate selection. For SBP, the candidates are ridge- or spline-based signed-residual corrections, shrunk and clipped using inner-fold choices; for DBP, the default candidate is the unchanged

base predictor. The exploratory DBP router predicts whether $|\Delta_{\text{DBP}}|$ exceeds a training-fold-derived 80th-percentile cutoff. Its threshold is selected without outer-fold labels, and its output does not alter DBP: it only requests repeat measurement or recalibration. We report discrimination, calibration, warning precision and recall, and warning rate. Because these analyses followed inspection of earlier AuroraBP results, their confidence intervals quantify internal resampling stability rather than independent confirmation.

For the independent multi-wavelength mechanism analysis, visible and infrared pulse-rate estimates are weighted by within-window spectral concentration and autocorrelation quality. The tone-blind optical arm uses no MST information. The guarded arm permits an MST-dependent wavelength-weight exponent selected within four participant-disjoint inner folds. It activates only when the proposed nonzero exponent improves mean inner-fold MAE by at least 0.10 BPM and improves all inner folds; otherwise it returns exactly to tone-blind optical fusion. Evaluation uses 20 repeats of five-fold participant outer cross-validation. The supplied script averages repeated out-of-fold predictions for each window before calculating participant MAE; these are repeated-cross-fit ensemble estimates, not mean per-repeat errors. The audit’s fixed equal-weight comparator is $(h_V + h_I)/2$, requires no fitting, and uses exactly the same 630 windows. Paired differences are averaged within each of 108 participants and bootstrapped 10,000 times (seed 20260910); those additional comparisons are post hoc with marginal intervals.

For VitalDB, the median BP and median features of three early valid 20-s windows form the case-specific calibration anchor and are excluded from fitting and scoring; up to 24 later windows per case are scored. A calibration-plus-context control uses the anchor value, elapsed time, age, sex, and preoperative hypertension. The PPG arm adds changes in pulse rate, spectral concentration and entropy, waveform and derivative energy, distributional shape, beat-interval variability, and signal quality. After the initial 55-subject pilot, this stricter PPG-versus-calibration-plus-context comparison was frozen for the expanded test; it was not prospectively registered. Candidate estimators, regularization, clipping, and the context/PPG blend were selected using only the 55 development subjects. The selected models were then fitted on development data and applied once to 228 cases from 227 disjoint held-out subjects. Both endpoints selected a constant PPG weight of 0.5: the final prediction is the arithmetic mean of context and context-plus-PPG predictions. Signal quality is an input, not a window-dependent blend weight. This is a separately fitted implementation, not an unchanged AuroraBP model transferred to VitalDB. Reference quality control required plausible SBP/DBP, local arterial stability, and pulse pressure; it did not use prediction error. Contrasts were averaged within subject before a 10,000-resample participant bootstrap. The holdout was not used for scaling, model selection, clipping, or blend selection.

Table 2: Refitted Repeated Participant-Level Validation (653 Participants; Three-Reading Anchor)

Arm	SBP MAE [95% CI]	DBP MAE [95% CI]
Calibration retention	9.790	6.767
No added context	8.133 [7.824, 8.455]	6.265 [6.019, 6.526]
Optical quality only	8.110 [7.802, 8.431]	6.273 [6.025, 6.535]
Wearable-state context	7.697 [7.407, 8.008]	6.139 [5.899, 6.391]
Protocol-enriched upper bound	7.523 [7.243, 7.815]	5.743 [5.527, 5.977]
Primary gain [97.5% CI]	0.436 [0.305, 0.573]	0.126 [0.053, 0.200]
Holm-adjusted paired p	0.00002	0.00013

5 Experiments and Results

5.1 Primary Repeated Participant-Level Validation

Table 2 reports the refitted comparison on 653 participants and 7,879 strict-endpoint measurements. The wearable-state arm achieves 7.697/6.139 mmHg SBP/DBP. Its paired gains over no added context are 0.436 and 0.126 mmHg. The corresponding 97.5% bootstrap intervals are [0.305, 0.573] and [0.053, 0.200] mmHg. Holm-adjusted paired sign-flip p values are 0.00002 and 0.00013; paired t sensitivity tests also pass the two-endpoint Holm threshold. These conditional internal-validation results support the added context comparison. The optical-quality-only changes are negligible and directionally mixed: 0.022 mmHg lower SBP MAE but 0.008 mmHg higher DBP MAE than no added context, providing no meaningful evidence of incremental value from the supplied optical-quality field alone. Posture and activity are protocol annotations; deployment requires independent wearable estimates. The protocol-enriched arm additionally uses phase, duration, and pressure-quality metadata and is interpreted as an internal upper bound. As a direct consistency check, weighting the regenerated Primary-CV Fitzpatrick-band MAEs in Supplementary Table S8 by their participant counts (516/112/25) reproduces the Table 2 wearable-state estimates exactly before rounding: 7.6971776/6.1385779 mmHg.

5.2 Matched Feature-Family Ablation

Table 3 reports the newly run secondary ablation. Adding the standard PPG descriptors to calibration, time, and wearable context reduced participant MAE by 0.187 mmHg SBP (95% CI 0.088–0.286) and 0.180 mmHg DBP (0.096–0.266). Adding the nonlinear/attractor family to that standard arm produced only 0.025 mmHg SBP (−0.018 to 0.069) and 0.009 mmHg

Table 3: Secondary Matched AuroraBP Feature-Family Ablation (653 Participants)

Arm	SBP MAE [95% CI]	DBP MAE [95% CI]
Calibration, time, wearable context	7.920 [7.607, 8.256]	6.318 [6.067, 6.589]
+ Standard PPG descriptors	7.733 [7.441, 8.048]	6.137 [5.898, 6.388]
+ Nonlinear/attractor family	7.709 [7.418, 8.017]	6.128 [5.887, 6.384]
Attractor-family gain [95% CI]	0.025 [−0.018, 0.069]	0.009 [−0.033, 0.050]
Holm-adjusted paired p	0.532	0.662

Table 4: Refitted Protocol-Enriched Models by Calibration Dose (653 Participants)

Target	Readings	Retention	Model MAE [95% CI]	Gain vs. retention
SBP	1	10.547	8.257 [7.933, 8.602]	2.290 [2.051, 2.525]
SBP	2	10.037	7.750 [7.458, 8.055]	2.286 [2.066, 2.509]
SBP	3	9.790	7.523 [7.243, 7.815]	2.267 [2.044, 2.489]
DBP	1	7.213	6.280 [6.026, 6.548]	0.934 [0.739, 1.129]
DBP	2	6.900	5.873 [5.648, 6.112]	1.027 [0.833, 1.215]
DBP	3	6.767	5.743 [5.527, 5.977]	1.024 [0.832, 1.211]

DBP (−0.033 to 0.050) further reductions. The Holm-adjusted paired sign-flip p values for those latter contrasts were 0.532 and 0.662. Thus the matched analysis supports an incremental contribution from the declared standard PPG descriptors but does not establish an attractor-family advantage.

5.3 Calibration Dose

Table 4 reports freshly refitted protocol-enriched models on identical participants and future measurements for each calibration dose. Ranking and fitting are repeated within the same participant folds. These dose gains belong to the protocol-enriched arm, not the lead wearable-state arm. Three readings are one-session repeated calibration, not one-reading calibration.

5.4 Prior-Anchored Isolation Ladder: Where the 26.2% Figure Really Comes From

Table 5 decomposes the strongest raw number in this study – a 2.084 mmHg (26.2%) SBP reduction – into its components under a controlled, feature-set-fixed ladder. **Eighty-three percent of that gain is attributable to re-anchoring alone** (predicting from the most

Table 5: Historical Controlled Isolation Ladder, Prior-Anchored (SBP, common dev-only feature set, 168 holdout participants)

Step	MAE	Gain [95% CI]	Share
Session anchor (legacy)	7.956	—	—
→ Prior anchor	6.216	+1.739 [1.296, 2.197]	83%
→ Absolute-error loss	6.114	+0.102 [−0.028, 0.240]	5% [†]
→ +Context	5.795	+0.319 [0.189, 0.446]	15%
→ +HCA features	5.872	−0.077 [−0.183, 0.024]	−4% [†]
Total (session → full)		+2.084 [1.637, 2.533]	26.2%

[†] CI crosses zero: not established at 95%.

recent prior reading rather than the session calibration), which on an in-lab protocol can mean a reference only minutes old; this is a genuine tracking result, but it is not a cuffless-replacement result, and the two must not be reported interchangeably (Section 3.3).

5.5 Secondary Pigmentation-Intervention Audit

Table 6 compares four configurations, all built on the prior-anchored absolute-error+context base model of Table 5 (chosen because it is the strongest available base for isolating the marginal effect of skin-tone modeling): *reference* (no added Fitzpatrick terms; indirect dependence is possible), *skin-features-only* (continuous Fitzpatrick terms and interactions, no rebalancing), *skin-features+balance* (adds participant rebalancing across Fitzpatrick bands and group-specific shrinkage gates learned on development data), and *updated-skin-aware* (the full configuration). All four share participants, targets, holdout, base features, estimator, and hyperparameters. **Skin features alone show no established benefit** (5.770 vs. 5.795 mmHg SBP reference, not CI-distinguishable). Adding the combined rebalancing and group-gating intervention moves SBP MAE from 5.770 to 6.086 mmHg, a point difference of 0.316 mmHg. A paired interval for this specific incremental contrast is not supplied; the reported 0.291 mmHg interval instead compares the full intervention with the original reference arm. These experiments do not isolate rebalancing from group gating. In the signed contrast reported in Table 6, reference minus skin-aware is therefore −0.291 mmHg [−0.420, −0.170]. The combined intervention also widens the between-Fitzpatrick-band MAE gap, from 0.556 to 0.772 mmHg (SBP) and 0.721 to 0.940 mmHg (DBP). Per-band results suggest harm in the two better-represented bands; the Fitzpatrick V–VI comparison is indeterminate and rests on only six evaluation participants. These historical fixed-split results do not support adopting the tested combined intervention but do not establish general fairness properties.

The supplementary subgroup table now reports newly refitted historical-method esti-

Table 6: Historical Fixed-Split Skin-Tone Model Comparison (168 Evaluation Participants)

Configuration	SBP MAE	DBP MAE	vs. blind
Reference	5.795	4.486	—
Skin-features-only	5.770	4.477	Neutral
+ Demographic balance	6.086	4.712	Harm
Updated skin-aware (full)	6.086	4.712	Harm
Paired diff. (reference — full), mmHg [CI]; negative favors reference:			
SBP	−0.291 [−0.420, −0.170]		
DBP	−0.226 [−0.345, −0.116]		
Between-band gap, reference (SBP/DBP)	0.556 / 0.721 mmHg		
Between-band gap, skin-aware (SBP/DBP)	0.772 / 0.940 mmHg		

mates alongside the regenerated primary wearable-state results. The companion change record documents the differences from the superseded subgroup values.

5.5.1 Repeated-CV pigmentation-correction intervention

Table 7 reports the historical V28 strict-cohort repeated-CV ablation retained from the development record; it was not refitted in the V38 primary rerun. Neither overall interval establishes a benefit; equivalence was not tested. The V–VI SBP subgroup shows a CI-supported 0.222 mmHg worsening under correction, whereas the I–II DBP subgroup shows a small CI-supported improvement. These subgroup findings are heterogeneous and exploratory; they do not support a general skin-aware-model benefit. Adding explicit Fitzpatrick interactions to the corrected channel did not establish overall benefit: SBP $\Delta = 0.001$ mmHg (95% CI −0.042 to 0.041) and DBP $\Delta = -0.037$ mmHg (−0.076 to 0.002), relative to the raw model.

An exploratory signal diagnostic collapsed all repeated observations to one median per participant before inference. Fitzpatrick type was associated with lower raw-channel PPG quality ($\rho = -0.157$, 95% CI −0.234 to −0.079), greater normalized attractor dispersion

Table 7: Pigmentation-Corrected vs. Raw Channel in Repeated Participant CV

Target	Scope	n	Raw	Corrected	Δ MAE [95% CI]
SBP	Overall	653	7.524	7.517	−0.007 [−0.042, 0.029]
SBP	I–II	516	7.487	7.482	−0.005 [−0.041, 0.032]
SBP	III–IV	112	7.683	7.620	−0.064 [−0.168, 0.036]
SBP	V–VI	25	7.571	7.793	0.222 [0.043, 0.401]
DBP	Overall	653	5.743	5.713	−0.030 [−0.062, 0.003]
DBP	I–II	516	5.681	5.639	−0.043 [−0.079, −0.005]
DBP	III–IV	112	6.064	6.077	0.014 [−0.066, 0.104]
DBP	V–VI	25	5.578	5.602	0.024 [−0.084, 0.131]

Δ = corrected minus raw; negative values favor correction.

($\rho = 0.511$, 0.447 to 0.571), and lower recorded optical-quality score ($\rho = -0.174$, -0.253 to -0.092). Associations with optical noise ratio ($\rho = 0.032$, -0.044 to 0.112) and raw spectral entropy ($\rho = 0.032$, -0.046 to 0.106) did not exclude zero. These are Fitzpatrick-associated observed-signal differences, not measured-melanin effects; device gain, contact pressure, anatomy, and protocol may confound them.

Under the historical session-anchored C0-versus-C2 comparison, per-band SBP MAE decreases in Fitzpatrick I–II ($7.865 \rightarrow 7.169$), III–IV ($8.151 \rightarrow 7.276$), and nominally V–VI ($9.042 \rightarrow 6.280$). The V–VI estimate is based on only six participants and is reported as indicative only; it cannot support a fairness or parity claim.

5.6 Independent Quality-Fusion and Guarded-Correction Test

Table 8 reports pulse-rate results on 108 participants. Quality-weighted fusion achieved 3.690 BPM MAE versus 4.366 for visible and 5.037 for infrared alone. A post-hoc audit of the same stored channel estimates added an equal-weight comparator, with 3.985 BPM MAE. Quality-minus-comparator paired differences were -0.676 BPM (95% CI -1.210 to -0.159) versus visible, -1.347 (-1.872 to -0.880) versus infrared, and -0.294 (-0.567 to -0.069) versus equal weighting. These marginal intervals support an overall pulse-rate benefit in this dataset; they do not establish BP utility. For MST 7–10, the quality-minus-equal-weight difference was -0.402 BPM (-0.929 to 0.010), which does not establish subgroup benefit. The exploratory MST guard activated in 5% of outer folds and achieved 3.703 BPM; its corrected-minus-tone-blind difference was 0.013 (-0.010 to 0.053), with no established overall benefit.

Table 8: Independent Multi-Wavelength PPG Mechanism Test

Arm	Overall MAE (BPM)	MST 7–10 MAE
Visible only	4.366	5.434
Infrared only	5.037	5.165
Equal-weight fusion (post hoc)	3.985	4.587
Quality-weighted fusion	3.690	4.184
Guarded MST-aware fusion	3.703	4.265




Figure 1: Participant-balanced pulse-rate MAE with marginal 95% bootstrap intervals. The equal-weight comparison was added post hoc; quality weighting improves overall error relative to equal weighting, while guarded MST correction does not establish additional benefit.

5.7 VitalDB Context-Plus-PPG Frozen Holdout

The expanded analysis used 55 development subjects to freeze the estimators and fixed blend, followed by one evaluation on 228 cases from 227 disjoint held-out subjects and 5,472 post-calibration windows. Table 9 separates the large gain over retaining the early calibration value from the stricter incremental PPG contrast. Calibration plus time/demographic context accounted for most of the improvement. The frozen fixed 50/50 blend reduced SBP MAE from 15.827 to 15.375 mmHg ($\Delta = -0.451$, 95% CI -0.781 to -0.121) and DBP MAE from 9.189 to 8.833 mmHg ($\Delta = -0.357$, -0.512 to -0.194). It improved over context for 67.0% of subjects for SBP and 70.5% for DBP. For windows with ≥ 10 mmHg reference change, the incremental gains were 0.707 mmHg SBP (0.337–1.074) and 0.501 mmHg DBP (0.274–0.726). These results support a small incremental contribution of the PPG-augmented implementation within the VitalDB holdout, but absolute errors remain high

Table 9: VitalDB Context-Plus-PPG Frozen Holdout (227 Subjects)

Target	Retention	Context	Context+PPG	PPG−Context [95% CI]
SBP	25.291	15.827	15.375	−0.451 [−0.781, −0.121]
DBP	13.098	9.189	8.833	−0.357 [−0.512, −0.194]




Figure 2: VitalDB locked-holdout participant-balanced MAE. Calibration-plus-context explains most of the gain over retention; the frozen fixed context/PPG blend adds small CI-supported incremental improvements for both endpoints.

and the test is retrospective and intraoperative rather than a prospective wearable validation.

5.8 State-of-the-Art Comparison

Supplementary Table S9 compares representative cuffless-BP studies by dataset, participant separation, and calibration. Direct numerical ranking is inappropriate because populations, BP ranges, reference protocols, and calibration burdens differ; record-level leakage alone can materially change reported performance (González et al., 2023). The wearable-state MAE of 7.697/6.139 mmHg is higher than the best calibrated reports and its DBP MAE is higher than the calibration-free result reported by Samimi and Dajani (2023). AuroraBP provides repeated internal validation with participant separation, whereas VitalDB evaluates a separately developed implementation in an intraoperative holdout. Neither analysis establishes

wearable generalization or superiority over models evaluated under different protocols.

6 Discussion

The primary AuroraBP comparison supports adding wearable-state context to the shared calibrated model, although posture and activity remain protocol annotations rather than independently estimated wearable states. The protocol-enriched arm achieves lower error but has greater dependence on acquisition metadata. Calibration dose materially affects that arm’s performance. Because the common core already includes baseline BP, elapsed-time candidates, and a UCI prior, these results do not isolate PPG utility beyond a calibration-plus-context-only model on AuroraBP.

The independent studies answer different questions. The post-hoc equal-weight comparison supports quality-weighted pulse-rate fusion on the multi-wavelength dataset. VitalDB supports a small incremental BP contribution from a separately fitted context-plus-PPG implementation, combined with context through a fixed 50/50 blend. It does not test quality-dependent blending or transfer of AuroraBP-fitted parameters. The combined evidence supports calibrated tracking research but does not establish one universally generalizing architecture, an attractor-specific advantage, or clinical interchangeability. Pigmentation interventions remain secondary; small samples, possible indirect feature dependence, preclude fairness conclusions.

Clinical accuracy boundary. The method remains a research prototype for calibrated trend tracking. On the historical measurement-level standards analysis, SBP achieved BHS Grade C and failed the reported AAMI/ISO Criterion 1 check because error SD was 9.654 mmHg, above 8 mmHg. VitalDB absolute MAE was also high (15.375/8.833 mmHg). These failures cannot be repaired by emphasizing relative improvements, rejecting difficult observations after seeing their errors, or treating a calibrated tracking model as a validated cuff-replacement device. Existing data can support tail-error diagnostics and prospectively defined repeat-measurement rules, but not a claim of clinical interchangeability.

7 Limitations

Validation and development. AuroraBP is a controlled, single-site protocol with sessions shorter than one hour. Repeated participant-level cross-validation prevents direct participant overlap but cannot remove optimism from earlier choices made using this dataset. The independent Appel–Theart analysis concerns pulse extraction rather than BP, and VitalDB

evaluates a separately developed retrospective intraoperative model rather than transfer of the AuroraBP model or prospective wearable performance.

Calibration and context. The primary result requires three initial cuff readings. The one-reading protocol-enriched model has higher MAE (8.257/6.280 mmHg), and a matching wearable-state dose analysis is unavailable. Posture and activity are recorded annotations rather than outputs of a deployed wearable classifier; phase, duration, and pressure-quality variables in the protocol-enriched arm may encode study structure.

Pigmentation and fairness. The repeated-CV cohort includes only 25 Fitzpatrick V–VI participants, and the historical fixed holdout includes six. The auxiliary-source pigmentation ranks are synthetic simulation inputs. The tested correction is one heuristic, and its null or adverse results establish neither equivalence nor the impossibility of a better intervention. The pulse-fusion study measures pulse-rate error, while VitalDB has no pigmentation labels; neither validates equitable BP performance.

Inference, standards, and reproduction. Holm adjustment covers only the two retrospectively designated primary endpoints; other analyses use exploratory marginal intervals. Participant-level summaries are used because repeated measurements are correlated. The DBP warning router has only 4.3% recall and is not a safety mechanism. BHS and AAMI/ISO criteria were designed for absolute cuff-device validation and are included only as context. The rerun reproduces modeling and inference from hashed feature banks but does not repeat raw-waveform extraction. The matched ablation tests the combined nonlinear/attractor family and does not isolate any one descriptor.

8 Conclusion

This study provides retrospective evidence for context-augmented, session-calibrated BP tracking. Wearable context and standard PPG descriptors produced small incremental improvements, whereas the combined nonlinear/attractor family did not show a reliable additional benefit. Independent analyses supported quality-weighted pulse-rate fusion and a small PPG contribution within a separately developed VitalDB model. The results do not establish transfer of one model across datasets, equitable or prospective performance, standards compliance, or cuff replacement.

Code and Reproduction Availability

The companion reproducibility package is publicly available on GitHub:

<https://github.com/farouze/pc-avct-reproducibility>

It contains rerun and statistical scripts, the matched feature-family ablation, preserved V26 helper functions, pinned dependencies, SHA-256 input hashes, model-selection manifests, aggregate outputs, and sanitized development notebooks. Participant-level data, row predictions, and paired participant contrasts are excluded under the dataset access conditions and retained locally for audit. The versioned GitHub release is the software archive cited here; the dataset DOI below identifies data, not this implementation.

Data Availability

AuroraBP is available from its provider upon application through the Zenodo record cited in (Microsoft Research, 2022) (DOI: 10.5281/zenodo.19099166) and acceptance of the Data Use Agreement. The authors are not permitted to redistribute AuroraBP participant-level data. BIDMC and MIMIC-II are available through PhysioNet, UCI Cuff-Less BP through the UCI Machine Learning Repository, the multi-wavelength data through Mendeley Data (Appel and Theart, 2025), and VitalDB through PhysioNet (Lee and Jung, 2022). Access conditions and citations appear in Section 4.1. No directly identifying participant data are included in the manuscript or figures.

Ethics Statement

This study is a secondary computational analysis of previously collected, deidentified human data. The authors recruited no participants, performed no intervention, had no participant contact, received no direct identifiers, and made no attempt to reidentify participants. Morgan State University’s published IRB process states that submission is not required for research using deidentified data or publicly available datasets (Morgan State University, 2023). Accordingly, this analysis did not require review by the Morgan State University Institutional Review Board, and no approval number was issued. The analysis followed the ethical principles of the Declaration of Helsinki and applicable local requirements. Ethics approvals and written informed-consent procedures for the original data collections are reported by the respective dataset providers and source publications (Mieloszyk et al., 2022; Microsoft Research, 2022; Pimentel et al., 2017; Appel and Theart, 2025; Lee et al., 2022; Lee and Jung, 2022). No additional consent or consent for publication was required because the authors used only deidentified secondary data and report no identifiable participant information. This study was not a clinical trial.

Author Contributions

Farouk Ganiyu Adewumi: conceptualization, methodology, software, validation, formal analysis, data curation, visualization, and writing of the original draft. Timothy Oladunni: conceptualization, methodology, supervision, and review and editing of the manuscript. Both authors approved this revised manuscript and the public code release.

Competing Interests

The authors declare no competing financial or non-financial interests.

Funding

This work was supported in part by Morgan State University’s Office of Research and Graduate Programs.

Acknowledgment

We thank the participants, investigators, and data stewards who made the secondary datasets used in this study available.

References

- Raghda Al-Halawani, Meha Qassem, and Panicos A. Kyriacou. Analysis of the effect of skin pigmentation and oxygen saturation on monte carlo-simulated reflectance photoplethysmography signals. *Sensors*, 25(2):372, 2025. doi: 10.3390/s25020372.
- John Allen. Photoplethysmography and its application in clinical physiological measurement. *Physiological Measurement*, 28(3):R1–R39, 2007. doi: 10.1088/0967-3334/28/3/R01.
- Anneri Appel and Rensu Theart. Photoplethysmography device data: Hardware and participant study insights. Mendeley Data, 2025.
- Christoph Bandt and Bernd Pompe. Permutation entropy: A natural complexity measure for time series. *Physical Review Letters*, 88(17):174102, 2002. doi: 10.1103/PhysRevLett.88.174102.

- David Dasa and Philip Davies. Skin tone and diagnostic equity in contactless blood pressure screening: A prospective observational field evaluation of remote photoplethysmography in nigeria. *BMJ Open*, 16(6):e119311, 2026. doi: 10.1136/bmjopen-2026-119311.
- Felipe M. Dias, Diego A. C. Cardenas, Marcelo A. F. Toledo, Filipe A. C. Oliveira, Estela Ribeiro, Jose E. Krieger, and Marco A. Gutierrez. Exploring the limitations of blood pressure estimation using the photoplethysmography signal. *Physiological Measurement*, 46(4):045007, 2025. doi: 10.1088/1361-6579/adcb86.
- Jean-Pierre Eckmann, S. Oliffson Kamphorst, and David Ruelle. Recurrence plots of dynamical systems. *Europhysics Letters*, 4(9):973–977, 1987. doi: 10.1209/0295-5075/4/9/004.
- Mohamed Elgendi, Richard Fletcher, Yongbo Liang, Newton Howard, Nigel H. Lovell, Derek Abbott, Kenneth Lim, and Rabab Ward. The use of photoplethysmography for assessing hypertension. *npj Digital Medicine*, 2:60, 2019. doi: 10.1038/s41746-019-0136-7.
- Mohamed Elgendi, Fridolin Haugg, Richard Ribon Fletcher, John Allen, Hangsik Shin, Ayemen Alian, and Carlo Menon. Recommendations for evaluating photoplethysmography-based algorithms for blood pressure assessment. *Communications Medicine*, 4:140, 2024. doi: 10.1038/s43856-024-00555-2.
- Ary L. Goldberger, Luis A. N. Amaral, Leon Glass, Jeffrey M. Hausdorff, Plamen Ch. Ivanov, Roger G. Mark, Joseph E. Mietus, George B. Moody, Chung-Kang Peng, and H. Eugene Stanley. PhysioBank, PhysioToolkit, and PhysioNet: Components of a new research resource for complex physiologic signals. *Circulation*, 101(23):e215–e220, 2000. doi: 10.1161/01.CIR.101.23.e215.
- Sergio González, Wan-Ting Hsieh, and Trista Pei-Chun Chen. A benchmark for machine-learning based non-invasive blood pressure estimation using photoplethysmogram. *Scientific Data*, 10:149, 2023. doi: 10.1038/s41597-023-02020-6.
- Sture Holm. A simple sequentially rejective multiple test procedure. *Scandinavian Journal of Statistics*, 6(2):65–70, 1979. doi: 10.2307/4615733.
- Carina Hörandtner, Martin Bachler, Walter Sehnert, Ines Mikisek, Thomas Mengden, Siegfried Wassertheurer, and Christopher C. Mayer. Attractor reconstruction for quantifying the arterial pulse wave morphology during device-guided slow breathing. *Cardiovascular Engineering and Technology*, 13(6):939–949, 2022. doi: 10.1007/s13239-022-00628-0.
- Rutendo Jakachira, Wenyan Yan, Sian C. Thomas, Yareli Macias-Sanchez, Lovisa Werner, Joshua A. Burrow, Shira Dunsiger, and Kimani C. Toussaint, Jr. Evaluation of a

- polarization-sensitive, dual-wavelength wearable photoplethysmography sensor across a range of skin tones. *Biophotonics Discovery*, 3(1):012509, 2026. doi: 10.1117/1.BIOS.3.1.012509.
- Mohamad Kachuee, Mohammad Mahdi Kiani, Hoda Mohammadzade, and Mahdi Shabany. Cuff-less high-accuracy calibration-free blood pressure estimation using pulse transit time. In *2015 IEEE International Symposium on Circuits and Systems (ISCAS)*, pages 1006–1009, 2015a. doi: 10.1109/ISCAS.2015.7168806.
- Mohamad Kachuee, Mohammad Mahdi Kiani, Hoda Mohammadzade, and Mahdi Shabany. Cuff-Less Blood Pressure Estimation. UCI Machine Learning Repository, 2015b.
- Jason Kostrna, Ekaterina Oparina, Cristina Palacios, Andres J. Rodriguez, JunZhu Pei, Ajmal Ajmal, and Jessica C. Ramella-Roman. PPG-based heart rate accuracy in diverse populations: Investigating inequities across body composition and skin tones. *arXiv preprint arXiv:2601.22377*, 2026. doi: 10.48550/arXiv.2601.22377.
- Hyung-Chul Lee and Chul-Woo Jung. VitalDB, a high-fidelity multi-parameter vital signs database in surgical patients. PhysioNet, 2022. Version 1.0.0.
- Hyung-Chul Lee, Yoonsang Park, Soo Bin Yoon, Seong Mi Yang, Dongnyeok Park, and Chul-Woo Jung. VitalDB, a high-fidelity multi-parameter vital signs database in surgical patients. *Scientific Data*, 9:279, 2022. doi: 10.1038/s41597-022-01411-5.
- Connor MacIsaac, Macros Nguyen, Alexander Uy, Tianmin Kong, and Ava Hedayatipour. A programmable gain calibration method to mitigate skin tone bias in PPG sensors. *Biosensors*, 15(7):423, 2025. doi: 10.3390/bios15070423.
- Norbert Marwan, M. Carmen Romano, Marco Thiel, and Jürgen Kurths. Recurrence plots for the analysis of complex systems. *Physics Reports*, 438(5–6):237–329, 2007. doi: 10.1016/j.physrep.2006.11.001.
- Suril Mehta, Nipun Kwatra, Mohit Jain, and Daniel McDuff. Examining the challenges of blood pressure estimation via photoplethysmogram. *Scientific Reports*, 14:18318, 2024. doi: 10.1038/s41598-024-68862-1.
- Microsoft Research. Aurora-BP Study Dataset. Zenodo, 2022. Available upon application under the provider’s Data Use Agreement.
- Rebecca Mieloszyk, Hope Twede, Jonathan Lester, Jeremiah Wander, Sumit Basu, Gabe Cohn, Greg Smith, Dan Morris, Sidhant Gupta, Desney Tan, Nicolas Villar, Moni

- Wolf, Sailaja Malladi, Matt Mickelson, Lauren Ryan, Lindsey Kim, Jeffrey Kepple, Susanne Kirchner, Emma Wampler, Riena Terada, Joel Robinson, Ron Paulsen, and T. Scott Saponas. A comparison of wearable tonometry, photoplethysmography, and electrocardiography for cuffless measurement of blood pressure in an ambulatory setting. *IEEE Journal of Biomedical and Health Informatics*, 26(7):2864–2875, 2022. doi: 10.1109/JBHI.2022.3153259.
- Morgan State University. IRB Process Synopsis. Office of Research Administration, 2023. URL https://www.morgan.edu/Documents/ADMINISTRATION/OFFICES/ora/IRB%20Forms/2023-03-21_IRB_Process_Synopsis.pdf. Accessed: 2026-09-10.
- Ramakrishna Mukkamala, Jin-Oh Hahn, Omer T. Inan, Lalit K. Mestha, Chang-Sei Kim, Hakan Töreyn, and Survi Kyal. Toward ubiquitous blood pressure monitoring via pulse transit time: Theory and practice. *IEEE Transactions on Biomedical Engineering*, 62(8): 1879–1901, 2015. doi: 10.1109/TBME.2015.2441951.
- Alexander Neshitov, Konstantin Tyapochkin, Evgeniya Smorodnikova, and Pavel Pravdin. Wavelet analysis and self-similarity of photoplethysmography signals for HRV estimation and quality assessment. *Sensors*, 21(20):6798, 2021. doi: 10.3390/s21206798.
- Timothy Oladunni and Farouk Ganiyu Adewumi. Attractor domain theory: A mathematical framework for cardiovascular attractor analysis with wearable photoplethysmography (PPG) validation. *arXiv preprint arXiv:2606.22039*, 2026a. doi: 10.48550/arXiv.2606.22039.
- Timothy Oladunni and Farouk Ganiyu Adewumi. Attractor-vascular coupling theory: Formal grounding and empirical validation for AAMI-standard cuffless blood pressure estimation from smartphone photoplethysmography. *arXiv preprint arXiv:2605.10871*, 2026b. doi: 10.48550/arXiv.2605.10871.
- Timothy Oladunni and Farouk Ganiyu Adewumi. Cardiac stability theory: An axiomatically grounded framework for continuous cardiac health monitoring via smartphone photoplethysmography. *arXiv preprint arXiv:2604.23876*, 2026c. doi: 10.48550/arXiv.2604.23876.
- Laura Osorio-Sanchez, James M. May, and Panicos A. Kyriacou. Evaluation of skin pigmentation effect on photoplethysmography signals using a vascular finger phantom with tunable optical and mechanical properties. *Journal of Biomedical Optics*, 30(11):117002, 2025. doi: 10.1117/1.JBO.30.11.117002.

- Marco A. F. Pimentel, Alistair E. W. Johnson, Peter H. Charlton, Douglas Birrenkott, Peter J. Watkinson, Lionel Tarassenko, and David A. Clifton. Toward a robust estimation of respiratory rate from pulse oximeters. *IEEE Transactions on Biomedical Engineering*, 64(8):1914–1923, 2017. doi: 10.1109/TBME.2016.2613124.
- Marco A. F. Pimentel, Alistair E. W. Johnson, Peter H. Charlton, and David A. Clifton. BIDMC PPG and Respiration Dataset. PhysioNet, 2018. Version 1.0.0.
- Joshua S. Richman and J. Randall Moorman. Physiological time-series analysis using approximate entropy and sample entropy. *American Journal of Physiology–Heart and Circulatory Physiology*, 278(6):H2039–H2049, 2000. doi: 10.1152/ajpheart.2000.278.6.H2039.
- Mohammed Saeed, Mauricio Villarroel, Andrew T. Reisner, Gari Clifford, Li-Wei Lehman, George Moody, Thomas Heldt, Tin Kyaw, Benjamin Moody, and Roger Mark. MIMIC-II Clinical Database. PhysioNet, 2011a. Version 2.6.0.
- Mohammed Saeed, Mauricio Villarroel, Andrew T. Reisner, Gari Clifford, Li-Wei Lehman, George Moody, Thomas Heldt, Tin H. Kyaw, Benjamin Moody, and Roger G. Mark. Multiparameter intelligent monitoring in intensive care ii (MIMIC-II): A public-access intensive care unit database. *Critical Care Medicine*, 39(5):952–960, 2011b. doi: 10.1097/CCM.0b013e31820a92c6.
- Hamed Samimi and Hilmi R. Dajani. A PPG-based calibration-free cuffless blood pressure estimation method using cardiovascular dynamics. *Sensors*, 23(8):4145, 2023. doi: 10.3390/s23084145.
- Ivandro Sanches, Victor V. Gomes, Carlos Caetano, Lizeth S. B. Cabrera, Vinicius H. Cene, Thomas Beltrame, Wonkyu Lee, Sanghyun Baek, and Otávio A. B. Penatti. MIMIC-BP: A curated dataset for blood pressure estimation. *Scientific Data*, 11:1233, 2024. doi: 10.1038/s41597-024-04041-1.
- Sanidhya Singh, Miles Romney Bennett, Chen Chen, Sooyoon Shin, Hamid Ghanbari, and Benjamin W. Nelson. Impact of skin pigmentation on pulse oximetry blood oxygenation and wearable pulse rate accuracy: Systematic review and meta-analysis. *Journal of Medical Internet Research*, 26:e62769, 2024. doi: 10.2196/62769.
- Floris Takens. Detecting strange attractors in turbulence. In *Dynamical Systems and Turbulence, Warwick 1980*, volume 898 of *Lecture Notes in Mathematics*, pages 366–381. Springer, 1981. doi: 10.1007/BFb0091924.

Ching-Fu Wang, Ting-Yun Wang, Pei-Hsin Kuo, Han-Lin Wang, Shih-Zhang Li, Chia-Ming Lin, Shih-Chieh Chan, Tzu-Yu Liu, Yu-Chun Lo, Sheng-Huang Lin, and You-Yin Chen. Upper-arm photoplethysmographic sensor with one-time calibration for long-term blood pressure monitoring. *Biosensors*, 13(3):321, 2023. doi: 10.3390/bios13030321.

Supplementary material

Farouk Ganiyu Adewumi and Timothy Oladunni

Description. Implementation details, historical analyses, negative results, clinical-standards context, subgroup diagnostics, and a comparison of representative cuffless blood-pressure studies.

Primary claims and conclusions are reported in the main manuscript.

Appendix A Supplementary Implementation Details

A. Multi-Wavelength Pulse Processing

For each device, visible and infrared pulse rates are estimated around the scheduled reference time using a window of up to 12 s, with at least 7.5 s and 150 samples required. Timestamps are converted to seconds, sorted, and deduplicated; signals are interpolated to a uniform grid using the median interval, with inferred sampling rate clipped to 20–500 Hz. A third-order 0.55–4 Hz Butterworth filter is applied forward and backward after detrending. Welch spectra use up to 8 s segments (at least 256 samples when available), with 50% overlap. The dominant frequency is sought in 0.65–3.5 Hz; spectral concentration c is the fraction of pulse-band power within 0.18 Hz of that peak.

An autocorrelation pulse estimate searches lags corresponding to 39–210 BPM. Let a be the selected autocorrelation peak clipped to $[0, 1]$, and h_F, h_A the spectral and autocorrelation rates. When $|h_F - h_A| \leq 15$ BPM, the channel estimate is $(ch_F + ah_A)/(c + a)$; otherwise the estimate with the larger reliability component is used. Channel quality is

$$q = \text{clip}(\sqrt{ca}g, 10^{-4}, 1), \quad (1)$$

$$g = \begin{cases} 1, & |h_F - h_A| \leq 15, \\ e^{-|h_F - h_A|/30}, & \text{otherwise.} \end{cases} \quad (2)$$

The tone-blind fusion estimate is $(q_V h_V + q_I h_I)/(q_V + q_I)$. The optional tone-dependent arm replaces the weights by $q_V e^{-\alpha z}$ and $q_I e^{\alpha z}$, with $z = \text{clip}[(\text{MST} - 1)/9, 0, 1]$. Candidate α values are $\{0, \pm 0.25, \pm 0.5, \pm 0.75, \pm 1, \pm 1.5, \pm 2, \pm 3\}$. Four participant folds within each outer training set score these fixed candidate rules. The selected nonzero exponent must improve mean fold MAE by at least 0.10 BPM and improve every inner fold; otherwise $\alpha = 0$. The seed is 20260827, with repeat/fold offsets defined in the companion script. These methods estimate pulse rate, not BP.

B. VitalDB Extraction and Model Selection

Eligible cases have all three required tracks and a case duration of at least 1,800 s. Cases are deterministically shuffled with seed 20260827; positions 1–60 define development and 61–300 define the candidate holdout. The saved audit contains 55 development cases/subjects and 228 eligible holdout cases from 227 subjects, with no subject overlap. Eligibility and sampling are retrospective, and selecting windows across the completed record is not an online scheduling policy.

Candidate centers start no earlier than 120 s, proceed in 30-s increments, and require at least 900 s of usable overlapping record. PPG is retained at approximately 100 Hz. A 20-s window must contain at least 75% of the expected samples and 85% finite values; gaps are interpolated. Values are clipped to the median plus/minus eight robust SD estimates, detrended, and filtered with a third-order 0.5–8 Hz bandpass (upper cutoff limited by sampling rate). Pulse quality is spectral concentration multiplied by $\exp[-\min(\text{interval CV}, 2)]$; windows below 0.04 are excluded. Pulse frequency search, spectral concentration, and beat-interval limits follow the companion extraction code. This quality filter is held fixed, but a no-quality-filter ablation is not reported.

Reference BP is the median of at least three plausible values within ± 5 s of the center. SBP must lie in 70–220 mmHg and DBP in 35–130 mmHg; local robust SD must not exceed 8 and 6 mmHg, respectively. Pulse pressure must lie in 20–130 mmHg. These reference-based exclusions do not use prediction errors but restrict the evaluated BP distribution. Cases require at least eight valid windows. The first three define median BP and feature anchors; up to 24 later windows are selected using evenly spaced indices over the valid future windows.

The context model uses anchor BP, $\log(1 + \text{elapsed seconds})$, age, sex, and preoperative hypertension (missing hypertension is coded as zero). The PPG model adds current and calibration-relative optical features listed by the columns function in the companion script. Candidate estimators comprise standardized ridge regression with 13 logarithmically spaced penalties from 10^{-3} to 10^3 and three absolute-loss gradient-boosting configurations. Development-only five-fold participant validation selects estimators, followed by a fixed blend weight from $\{0, 0.25, 0.5, 0.75, 1\}$; a nonzero weight requires at least 0.10 mmHg improvement over context. Development selection is not independently nested across the estimator and blend choices, but holdout observations are excluded from both. Delta predictions are clipped at training-target 1st/99th percentiles.

Both endpoints selected ridge context models with penalty approximately 31.6228. The SBP PPG model uses absolute-loss boosting, learning rate 0.03, 15 leaves, regularization 5, 250 iterations, and minimum leaf size 20; the DBP PPG model uses ridge with penalty approximately 31.6228. Both selected a constant 0.5 PPG weight. The saved predictions satisfy this fixed-blend identity to numerical precision. Errors are averaged over all windows from all cases of a subject before subject bootstrap; the subject contributing two cases remains one bootstrap unit.

C. Audit Scope and Reproduction Status

The independent-dataset saved-result audit uses 10,000 participant bootstrap resamples with seed 20260910 for the pulse comparisons and verifies window uniqueness, case/subject counts, development/holdout separation, the fixed VitalDB blend, and the pulse-fusion equation. The AuroraBP analysis was separately rerun: the context arms and all three calibration doses were refitted across three repeats of five participant folds. The package records hashed feature inputs, software versions, selected features and gates, row-level outer-fold predictions, and paired participant contrasts. The historical subgroup model was also refitted. The rerun starts from feature banks; raw-waveform extraction is outside its scope.

Appendix B

Historical Development and Auxiliary Analyses

The following analyses preserve the development record. They are secondary or exploratory. The final subgroup subsection reports newly refitted estimates; development codes C0–C4 refer to the historical arms.

A. Exploratory Endpoint-Specific Calibration and Warning Router

Table S2 reports the historical post-primary endpoint analysis retained from the development record; it was not refitted in the V38 primary rerun. A nested, tightly clipped residual calibrator reduced SBP participant-level MAE from 7.611 to 7.555 mmHg, a paired improvement of 0.056 mmHg (95% CI 0.020–0.091; 55.9% of participants helped). The absolute gain is small, and tail behavior changed little (P90 absolute error 16.862 $\rightarrow$ 16.820 mmHg; P95 22.184 $\rightarrow$ 22.120 mmHg). DBP candidate specialization did not improve inner-fold performance consistently, so the base prediction was retained at 5.784 mmHg. This endpoint-specific policy is more conservative than adopting one post-processing rule for both targets. The exploratory DBP router was therefore frozen as warning-only. Extreme DBP change prevalence was 20.5% (95% CI 19.1–21.9%). The router achieved ROC AUC 0.794 (0.778–0.809), average precision 0.517 (0.483–0.555), and Brier score 0.130 (0.122–0.137). At the nested operating threshold, warning precision was 0.773 (0.695–0.854), but recall was only 0.043 (0.035–0.052) and the warning rate was 1.1% (0.9–1.4%). The router is thus a rare high-precision flag at the evaluated threshold; it is not a screen

for most extreme changes, does not reduce reported DBP MAE, and must not be interpreted as a clinical safety mechanism.

B. Historical Duplicate-Reading Noise Diagnostic

Table S3 retains a historical duplicate-reading diagnostic, not a universal achievable-error bound. Estimating per-reading noise as $\sigma_{P-S}/\sqrt{2}$ assumes independent, equal-variance observer errors. The supplied code distinguishes a stored average of two readings from a single-reading target and propagates target and anchor errors for a legacy two-time-point difference. Its 1.754/2.412 mmHg Gaussian absolute-noise estimates and 2.6%/9.2% variance ratios therefore must not be treated as validated floors for the primary three-reading anchor. Shared observer error, physiological change during measurement, and the actual target-averaging rule can change the estimates. A primary-endpoint noise calculation would require those components and their covariances; subtracting a putative floor from model MAE does not measure recoverable physiological error.

C. Recalibration Cadence

The historical code evaluated nominal recalibration intervals of 1, 3, 6, 12, and 24 h, with nearly identical MAE (8.593–8.594 mmHg SBP; 6.020–6.022 mmHg DBP). Those labels do not establish distinct effective recalibration schedules in the analyzed short laboratory sessions. The stored comparison lacks an event-count and elapsed-time audit sufficient to interpret a cadence effect. We therefore withdraw the inference that cadence beyond one hour is immaterial and make no recommendation about recalibration frequency from this analysis.

D. Multi-Window Feature Averaging

Averaging attractor features over multiple 10 s windows per recording (up to 3 available) was hypothesized to reduce measurement noise (Table S4 shows within/between SD ratios up to approximately 0.45). It did not help: holdout MAE with multi-window averaging (6.393/4.828 mmHg SBP/DBP) was statistically indistinguishable from, and nominally slightly worse than, a matched single-window configuration (6.346/4.803 mmHg), both close to the legacy single-window baseline (6.315/4.789 mmHg). We report this as a negative result (Table S1).

E. Historical Fixed-Split Session-Anchored Ladder

During earlier development, the model-development ladder was evaluated on a fixed 501/168 participant split. Table S5 preserves that analysis for transparency: C4 reduced MAE from 8.856/6.582 to 7.147/5.579 mmHg. Because the same 168 participants were examined across successive ablations, these values are historical development evidence and are not the paper’s primary validation result. The corresponding repeated participant-validation

TABLE S1
Interventions Tested and Not Adopted (Elimination Log)

Intervention	Times tested	Outcome
Absolute-error loss (vs. squared)	3	Never CI-supported: +0.152 [−0.001, 0.311] (51.8% helped); +0.102 [−0.028, 0.240]; +0.163 [−0.008, 0.345] (exactly 50.0% helped).
HCA feature augmentation	2	No CI-supported benefit: −0.077 [−0.183, 0.024]; +0.043 [−0.026, 0.115] (53.0% helped).
Signal-level pigmentation (Beer–Lambert) correction	2	First variant increased error; the repeated-CV corrected channel changed MAE by −0.007 [−0.042, 0.029] SBP and −0.030 [−0.062, 0.003] DBP – no established overall benefit (the main manuscript’s repeated-CV pigmentation-correction subsection).
Demographic (Fitzpatrick-band) rebalancing / group gating	5	Consistently harmful: e.g., −0.291 [−0.420, −0.170] mmHg SBP; widens between-band gap in every run (the main manuscript’s Secondary Pigmentation-Intervention Audit).
Multi-window feature averaging	1	No measurable benefit; nominally worse (Section B-D).
Nested gate/threshold search (early pipeline)	1	Higher error than a plainer reference baseline (8.253/8.183 vs. 8.109 mmHg SBP).
Passed sanity check: metadata-only permutation control	1	AUC = 0.492 ($\approx$ chance) – is consistent with chance under this permutation control; does not exclude other confounds.

TABLE S2
Exploratory Endpoint-Specific Nested Analysis (653 Participants)

Target	Arm	MAE	P90 AE	P95 AE
SBP	Base	7.611	16.862	22.184
SBP	Clipped residual	7.555	16.820	22.120
DBP	Base	5.784	12.043	15.630
DBP	Selected policy	5.784	12.043	15.630
SBP paired gain: 0.056 mmHg [0.020, 0.091]; DBP gain: 0.000.				

TABLE S3
Historical Duplicate-Reading Diagnostic (8,918/8,916 SBP/DBP pairs)

Quantity (mmHg)	SBP	DBP
σ (primary – secondary reading)	3.108	4.275
σ_{cuff} (per-reading noise)	2.198	3.023
Legacy Gaussian absolute-noise estimate	1.754	2.412
Observed SD of Δ	13.70	9.98
Legacy noise/observed variance ratio	2.6%	9.2%
Residual SD under legacy assumptions	13.52	9.51
Early reference model MAE	8.11	–

estimates in the main manuscript’s primary validation table are modestly worse (7.523/5.743 mmHg for protocol-enriched context), consistent with mild optimism from the fixed split rather than a collapse of the effect.

F. Calibration-Row Leakage Audit

Table S6 reports all four checks from the main manuscript’s calibration-row leakage audit. Check A shows calibration-flagged rows are smaller in mean $|\Delta|$ than unflagged rows (6.774 vs. 9.902 mmHg SBP, development set) but not trivially near-zero, and the fraction of rows within 1 mmHg is nearly identical between flagged and unflagged rows (9.8% vs. 8.6%). The decisive Check D – flag removed from features, flagged rows excluded from both fit and scoring – yields a context gain of +0.614 mmHg SBP [0.346, 0.889] and +0.513 mmHg

TABLE S4
Historical Within- and Between-Recording Feature SD

Feature	Within SD	Between SD	SD ratio
σ_m	0.0046	0.0102	0.449
CSI	0.0043	0.0138	0.311
Recurrence rate	0.0040	0.0130	0.305
Kurtosis	0.0902	0.3086	0.292
RQA entropy	0.0734	0.2520	0.291
Skewness	0.0386	0.1389	0.278
Sample entropy	0.0133	0.0527	0.252
Determinism	0.0114	0.0475	0.240
Spectral entropy	0.0092	0.0488	0.189
Permutation entropy	0.0028	0.0190	0.150

TABLE S5
Historical Fixed-Split Session-Anchored Ladder (168 Evaluation Participants)

Arm	SBP MAE [95% CI]	DBP MAE [95% CI]
Calibration retention	8.856	6.582
C0 legacy (squared)	7.956 [7.424, 8.501]	6.120 [5.641, 6.647]
C1 absolute loss	7.793 [7.270, 8.316]	6.014 [5.513, 6.546]
C2 +context	7.156 [6.740, 7.618]	5.651 [5.186, 6.146]
C3 +context+HCA	7.113 [6.703, 7.550]	5.636 [5.178, 6.126]
C4 +context (squared)	7.147 [6.736, 7.593]	5.579 [5.174, 6.004]
Reduction, calib. $\rightarrow$ C4	1.709 (19.3%)	1.003 (15.2%)

DBP [0.268, 0.743], both CI-supported and close to the unaudited historical figures. Leave-one-out attribution (Figure S2) shows larger predictive contributions from recording duration and post-exercise activity, not the calibration flag (LOO damage for the flag: 0.0000 SBP / 0.0027 DBP). One honest qualification: the single largest contributor to the context gain is recording duration (0.244 mmHg SBP / 0.366 mmHg DBP of leave-one-out damage), which is an acquisition property (likely proxying protocol step or signal stability) rather than a purely physiological state variable. “Context = posture and activity” is only partly right; post-exercise activity (0.154/0.144 mmHg) is the clearest genuinely physiological




Fig. S1. Historical fixed-split session-anchored ladder with 95% participant-bootstrap CIs. This figure documents model development; the main manuscript’s primary validation table provides the primary repeated participant-validation result.

TABLE S6
Calibration-Row Leakage Audit (168 holdout participants)

Check	SBP gain [CI]	DBP gain [CI]	Verdict
B (flagged excl.)	0.741 [0.440, 1.062]	0.606 [0.372, 0.854]	Supp.
C (flag dropped)	0.809 [0.501, 1.122]	0.538 [0.309, 0.770]	Supp.
D (strict)	0.614 [0.346, 0.889]	0.513 [0.268, 0.743]	Supp.

contributor, with posture third.

G. Clinical Standards Context

Table S7 reports BHS grading and AAMI/ISO 81060-2 Criterion 1 for the C4 session-anchored model. The model improves cumulative within-5/10/15 mmHg percentages over both calibration retention and the legacy baseline for both SBP and DBP, and reaches BHS Grade B for DBP (Grade C for SBP); it passes AAMI Criterion 1 for DBP (mean error 0.645, SD 7.992 $\leq$ 8 mmHg) but not for SBP (SD 9.654 $>$ 8 mmHg). We report both gradings as indicative context, not a validation claim: BHS and AAMI/ISO were written for cuff devices measuring absolute BP against an independent reference, and this model is calibrated per participant.

H. What Did Not Work

Table S1 summarizes every intervention tested and rejected in this study, alongside one passed sanity check. We regard this table as carrying as much evidentiary weight as Table S5: a result that survives repeated, adversarially-minded attempts to break it is more credible than one presented without them.

I. Refitted Fixed-Estimator and Primary-Cohort Subgroups

The historical prior-anchored estimator was refitted under the recorded environment, and Table S8 reports its newly generated predictions. These estimates supersede the earlier rounded subgroup values; they do not assert bitwise recovery of the original run. The companion change record retains the old-value comparison. The same estimator is applied within each analysis before




Fig. S2. Leave-one-out attribution of the session-anchored context gain (Check D). The largest contributor is recording duration – an acquisition property, not a purely physiological one – followed by post-exercise activity.

TABLE S7
BHS Grading and AAMI/ISO 81060-2 Criterion 1
(Session-Anchored C4 Model, Measurement-Level)

Target/Arm	$\leq 5\text{mm}$	$\leq 10\text{mm}$	$\leq 15\text{mm}$	Grade	AAMI-1
SBP calibration	45.9%	72.3%	83.9%	D	Fail
SBP baseline (C0)	45.8%	73.8%	86.6%	C	Fail
SBP model (C4)	48.0%	75.1%	89.7%	C	Fail
DBP calibration	56.0%	82.9%	92.2%	B	Fail
DBP baseline (C0)	56.0%	83.1%	93.6%	B	Fail
DBP model (C4)	58.7%	85.6%	94.9%	B	Pass

predictions are grouped by Fitzpatrick band. The primary wearable-state summary uses the independently regenerated repeated-CV predictions, not the historical fixed split. These are exploratory marginal intervals. Differences in sample size, anchoring, and validation design prevent direct comparison between the two analysis blocks. In particular, the small V–VI samples do not establish parity or equivalence, irrespective of the ordering of the point estimates.

References

- David Dasa and Philip Davies. Skin tone and diagnostic equity in contactless blood pressure screening: A prospective observational field evaluation of remote photoplethysmography in nigeria. *BMJ Open*, 16(6): e119311, 2026. doi: 10.1136/bmjopen-2026-119311.
- Sergio González, Wan-Ting Hsieh, and Trista Pei-Chun Chen. A benchmark for machine-learning based non-invasive blood pressure estimation using photoplethysmogram. *Scientific Data*, 10:149, 2023. doi: 10.1038/s41597-023-02020-6.
- Jingon Joung, Chul-Woo Jung, Hyung-Chul Lee, Moon-Jung Chae, Hae-Sung Kim, Jonghun Park, Won-Yong Shin, Changhyun Kim, Minhyung Lee, and Changwoo Choi. Continuous cuffless blood pressure monitoring using photoplethysmography-based PPG2BP-net for high intrasubject blood pressure variations. *Scientific Reports*, 13:8605, 2023. doi: 10.1038/s41598-023-35492-y.
- Timothy Oladunni and Farouk Ganiyu Adewumi. Attractor-vascular coupling theory: Formal grounding and empirical validation for AAMI-standard cuffless blood pressure estimation from smartphone photoplethysmography. *arXiv preprint arXiv:2605.10871*, 2026. doi: 10.48550/arXiv.2605.10871.

TABLE S8
Regenerated Participant MAE by Fitzpatrick Band

Analysis	Target	Band (n)	MAE	95% CI
Prior refit	SBP	I–II (133)	5.837	[5.513, 6.178]
Prior refit	SBP	III–IV (29)	5.568	[4.703, 6.549]
Prior refit	SBP	V–VI (6)	5.341	[3.847, 6.841]
Prior refit	DBP	I–II (133)	4.439	[4.118, 4.795]
Prior refit	DBP	III–IV (29)	4.781	[3.674, 6.370]
Prior refit	DBP	V–VI (6)	4.020	[2.931, 5.197]
Primary CV	DBP	I–II (516)	6.078	[5.805, 6.362]
Primary CV	DBP	III–IV (112)	6.435	[5.862, 7.045]
Primary CV	DBP	V–VI (25)	6.053	[5.065, 7.122]
Primary CV	SBP	I–II (516)	7.680	[7.355, 8.022]
Primary CV	SBP	III–IV (112)	7.786	[7.050, 8.610]
Primary CV	SBP	V–VI (25)	7.646	[6.052, 9.471]

TABLE S9
Representative Cuffless BP Studies

Method	Dataset (N)	Split	Calibration	SBP	DBP	Skin-tone report
PPG2BP-Net (Joung et al., 2023)	Intraoperative (4,185 participants)	Participant-independent	Required	$0.209 \pm 7.509^*$	$0.150 \pm 4.549^*$	Not reported
Calibration-free CV-dynamics (Samimi and Dajani, 2023)	UCI/MIMIC-II (225 participants)	Participant-independent	None	7.41	3.32	Not reported
Benchmark baseline (González et al., 2023)	Four public PPG–BP datasets	Participant-level CV	Mixed	15.60 (Sensors)	–	Not reported; quantifies leakage inflation
AVCT companion work (Oladunni and Adewumi, 2026)	Smartphone PPG	Single-point calibration	1 reading	2.05	1.67	Not reported
rPPG equity field study (Dasa and Davies, 2026)	Nigeria, Fitzpatrick V/VI (306 enrolled; 249 analyzed)	Independent cuff reference	None	15.4	10.9	Fitzpatrick V/VI
Wearable-state model (this study)	AuroraBP (653 participants)	3×5-fold participant CV	3 readings	7.697 [7.407, 8.008]	6.139 [5.899, 6.391]	Fitzpatrick I–VI; V–VI underpowered
Protocol-enriched upper bound	AuroraBP (653 participants)	Same internal CV	3 readings	7.523 [7.243, 7.815]	5.743 [5.527, 5.977]	Protocol metadata limit transportability
Context-plus-PPG holdout	VitalDB (227 held-out participants)	Locked participant holdout	3 windows	15.375	8.833	Not available; temporal test only

* Mean error $\pm$ SD rather than MAE; not directly comparable with the other rows.

Hamed Samimi and Hilmi R. Dajani. A PPG-based calibration-free cuffless blood pressure estimation method using cardiovascular dynamics. *Sensors*, 23(8): 4145, 2023. doi: 10.3390/s23084145.